



Facilitating SQL Query Composition and Analysis

Zainab Zolaktaf, Mostafa Milani, Rachel Pottinger

Department of Computer Science

University of British Columbia



Motivation: Facilitating SQL Query Composition and Analysis

- SQL query composition can be fundamentally difficult for users
 - Requires several cycles of tuning and execution of costly queries
- To write efficient SQL queries users can
 - Gain knowledge of database schema and tuples
 - Use [hints](#) or tutorials available on the system
 - E.g., On SDSS users are advised to write a ``Count`` query first!
- [Our goal](#): predict SQL query performance properties - prior to execution



Motivation: Facilitating SQL Query Composition and Analysis

A new SQL query

$Q_*, y_* = ?$

```
SELECT q.name AS qname,  
       dbo.fDistanceArcMinEq(q.ra,q.dec,p.ra,p.dec), ...  
FROM SpecObj AS s,  
     SDSSSQL010.MYDB_670681563.test.QSOQuery1_DR5 AS q, PhotoObj  
     AS p  
WHERE ((s.bestobjid=p.objid) AND (s.ra BETWEEN 185 AND 190) AND  
       ...) ORDER BY q.ra
```

Goal

Predict performance
properties of Q_* , prior to
submitting it the database


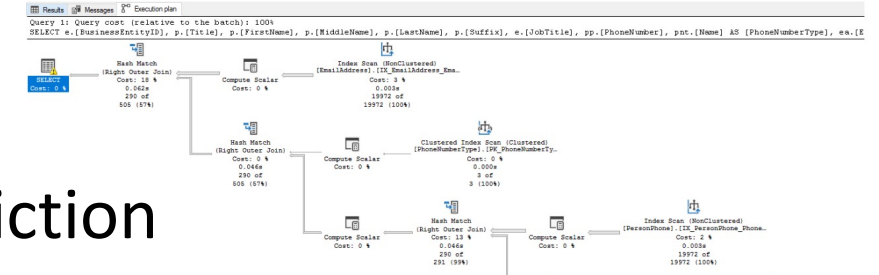
Output

- Answer size (y_*^a) = 304 rows
- CPU time (y_*^c) = 105.37 sec
- Error class (y_*^e) = success
- Session class (y_*^s) = browser

Goal
• Predict performance
properties of Q_* , prior to
submitting it the database

Challenges: Database Instance

- Existing models for query performance prediction
 - System side applications (e.g., admission control, query optimization [LKNC12])
- Use query execution plan
 - Need database instance and statistics
- Problems
 - Query execution plan can be imprecise [LGMB15]
 - Limited access to database instance?
 - Sources on the hidden web
 - Customers of cloud data warehouses
 - Spotify, HSBC use Google BigQuery



Output

- Cardinality estimates?
- Cost estimates?
- Error class (y_*^e) = success
- Session class (y_*^s) = browser



Challenges: Large-scale Query Workloads

$$W = \{(Q_i, y_i)\}_{i=1}^n$$

```
SELECT n.name AS nname
FROM S
WHERE j.outpr
SELECT j.target, cast(j.estimate AS varchar) AS queue,...
FROM S
WHERE j.outpr
SELECT p.objid, p.ra, p.dec, p.u,
      p.g, p.r, p.i, p.z
FROM PhotoObj AS p
WHERE type=6
AND p.ra BETWEEN (156.519031-0.200000)
AND (156.519031+0.200000)
AND p.dec BETWEEN (62.835405-0.200000)
AND (62.835405+0.200000)
ORDER BY p.objid
```

1. Sloan Digital Sky Server (SDSS) [RTS14]

- Scientific computing domain
- Extracted ~600K SQL queries

2. SQLShare [JMH16]

- SQL-as-a-Service platform
- Users upload data, write queries
- Contains ~27K SQL queries

- SQL Query workload (W)
 - Collection of labeled SQL queries submitted in the past
 - Labels are **actual observations**
 - Eliminate biases e.g., cardinality misestimates
 - Easily logged by DBMS
- Need **large-scale and real-world** query workloads
 - Reveal usage patterns from a variety of users



Problem Formulation: Facilitating SQL Query Composition and Analysis

Collection of labeled SQL queries

$$W = \{(Q_i, y_i)\}_{i=1}^n$$

```
SELECT n.name AS nname,
FROM S
WHERE j.outpr
SELECT j.target, cast(j.estimate AS varchar) AS queue, ...
FROM S
WHERE j.outpr
SELECT p.objid, p.ra, p.dec, p.u,
p.g, p.r, p.i, p.z
FROM PhotoObj AS p
WHERE type=6
AND p.ra BETWEEN (156.519031-0.200000)
AND (156.519031+0.200000)
AND p.dec BETWEEN (62.835405-0.200000)
AND (62.835405+0.200000)
ORDER BY p.objid
```

A new SQL query

$$Q_*, y_* = ?$$

```
SELECT q.name AS qname,
dbo.fDistanceArcMinEq(q.ra, q.dec, p.ra, p.dec), ...
FROM SpecObj AS s,
SDSSSQL010.MYDB_670681563.test.QSOQuery1_DR5 AS q, PhotoObj
AS p
WHERE ((s.bestobjid=p.objid) AND (s.ra BETWEEN 185 AND 190) AND
...) ORDER BY q.ra
```

Goal

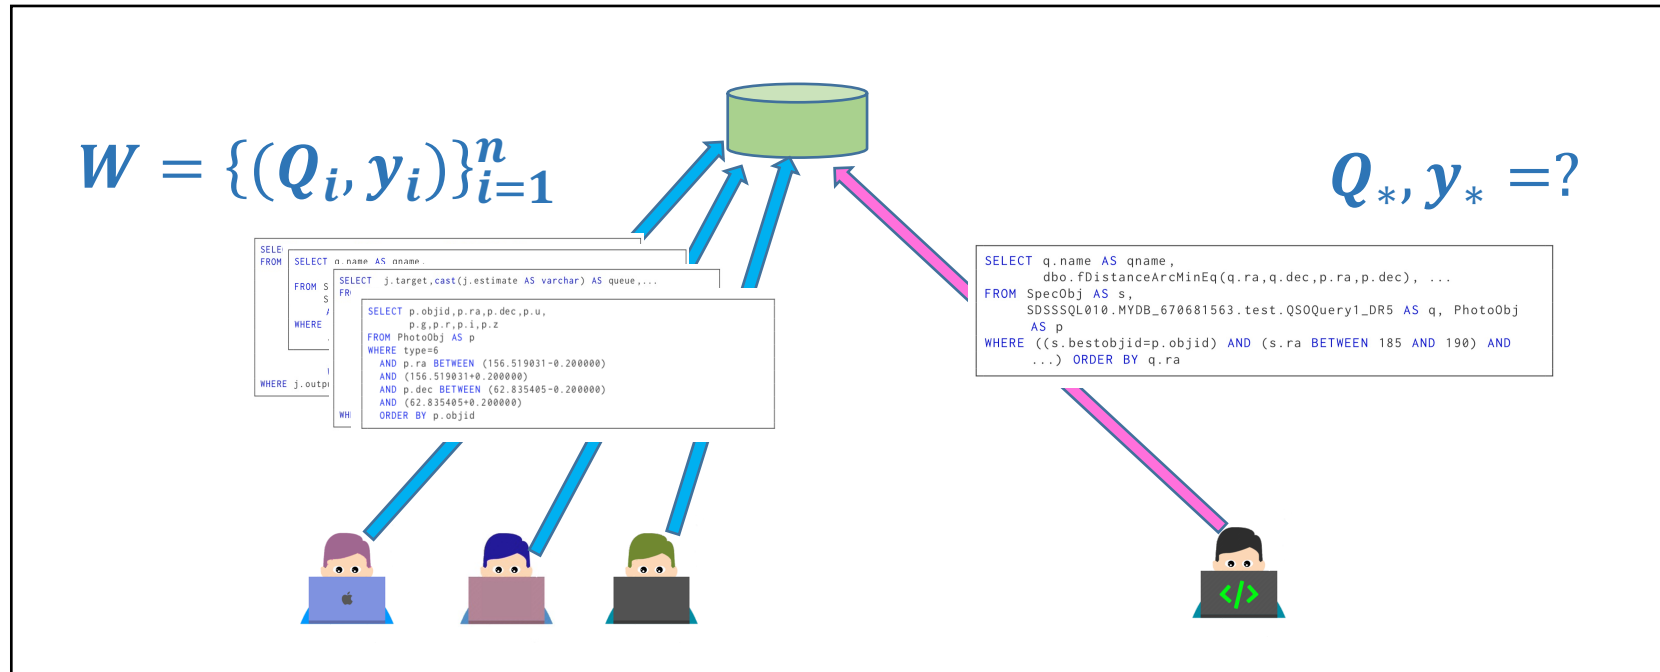
Predict performance properties of Q_* , prior to submitting it the database

Output

- Answer size (y_*^a) = 304 rows
- CPU time (y_*^c) = 105.37 sec
- Error class (y_*^e) = success
- Session class (y_*^s) = browser

Goal
• Predict performance properties of Q_* , prior to submitting it the database

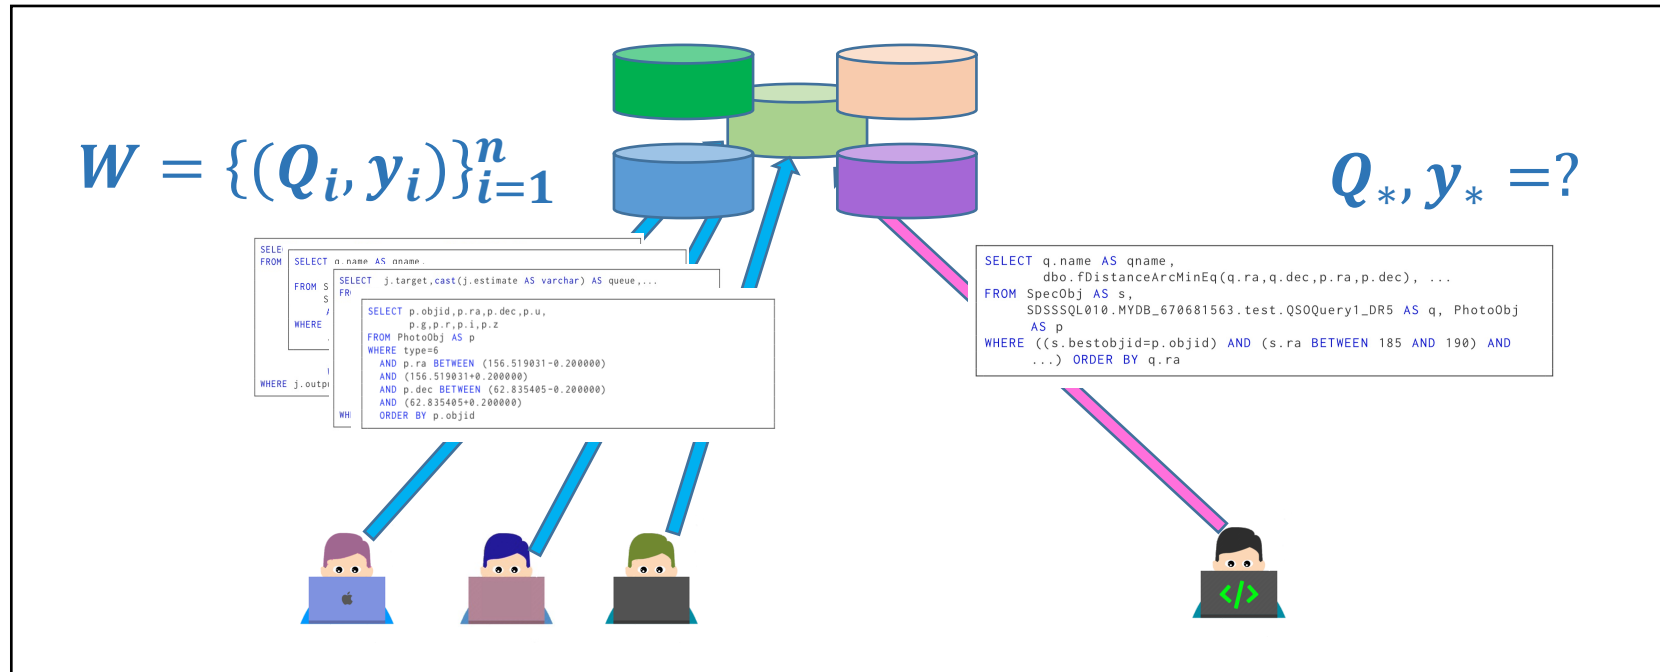
Approach Overview: Different Settings



1. **Homogeneous Instance:** Q_* and the queries in W are posed to the same database instance



Approach Overview: Different Settings

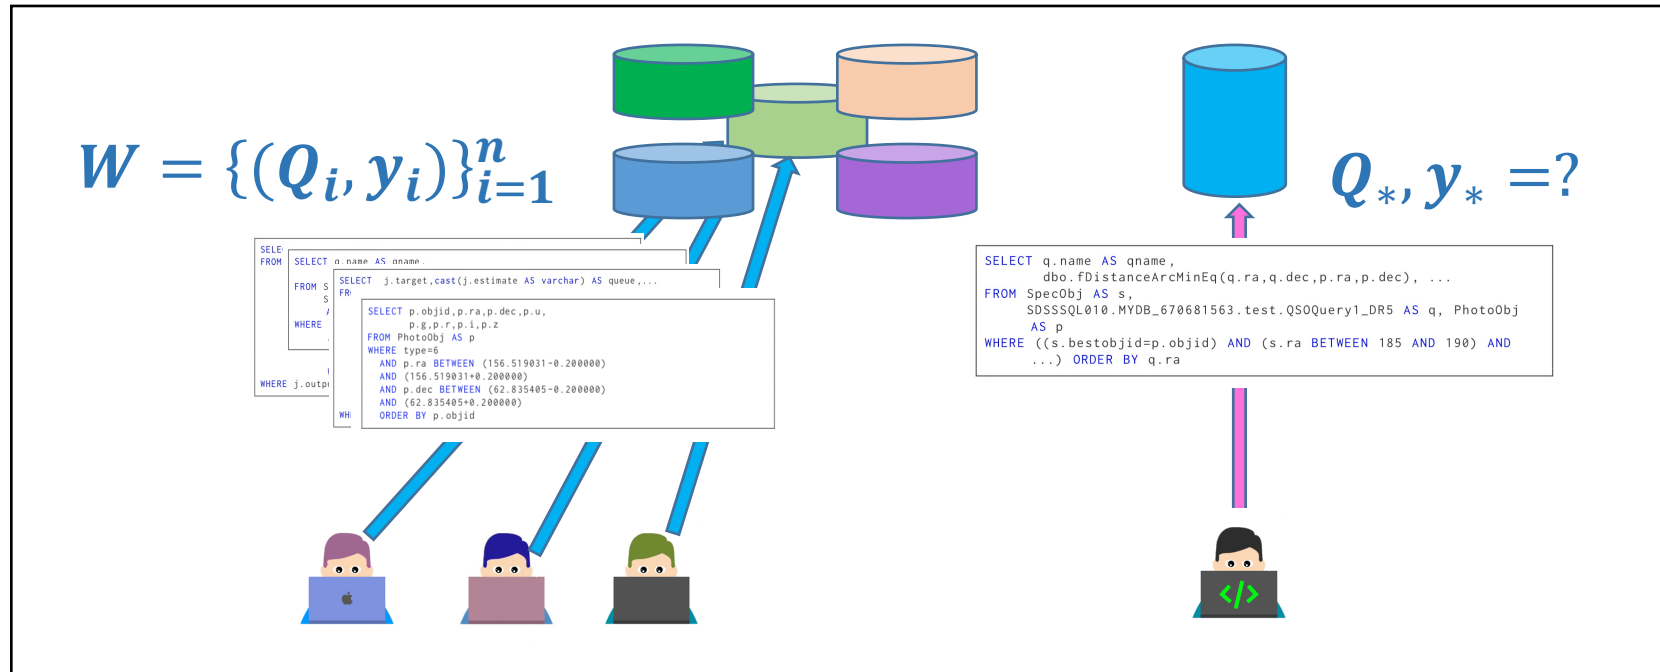


SQLShare

2. **Homogeneous Schema:** Q_* and the queries in W are posed to different database instances with the same schema in the same DBMS



Approach Overview: Different Settings



SQLShare




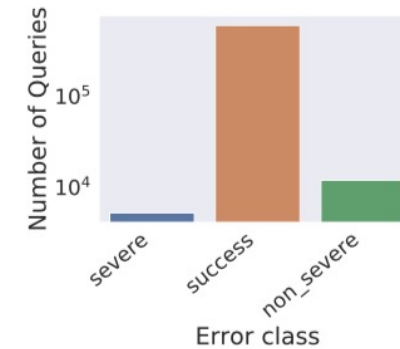
3. **Heterogeneous Schema:** Q_* and the queries in W are posed to different databases with different schemas that run in the same DBMS



Approach Overview: Workload Analysis



- Perform workload analysis for
 - Better **model selection**
 - Better **model evaluation**
- SQL query statements
 - Digits and mathematical equations in statements
 - Affect query performance, e.g., answer size
 - Range in complexity w.r.t. length, #joins
- SQL query labels
 - Classification labels are imbalanced
 - Regression labels had a wide range

```
SELECT q.name AS qname,  
       dbo.fDistanceArcMinEq(q.ra,q.dec,p.ra,p.dec), ...  
FROM SpecObj AS s,  
     SDSSSQL010.MYDB_670681563.test.QS0Query1_DR5 AS q, PhotoObj  
     AS p  
WHERE ((s.bestobjid=p.objid) AND (s.ra BETWEEN 185 AND 190) AND  
       ...) ORDER BY q.ra
```




Approach Overview: Models Evaluated

- To **establish baselines** we examined a broad set of models
 - Models that do not consider SQL query statement
 - Most frequent class (**mfreq**) classifier
 - **Median** of distribution for regression
 - Models that do consider SQL query statement
 - Query **statement representation?**
 - Bag-of-n-grams + **TFIDF**
 - Shallow Convolutional Neural Network (**CNN**)
 - 3-Layer Long Short-Term Memory (**LSTM**)
 - Applied at character and word level




Results: CPU Time Prediction

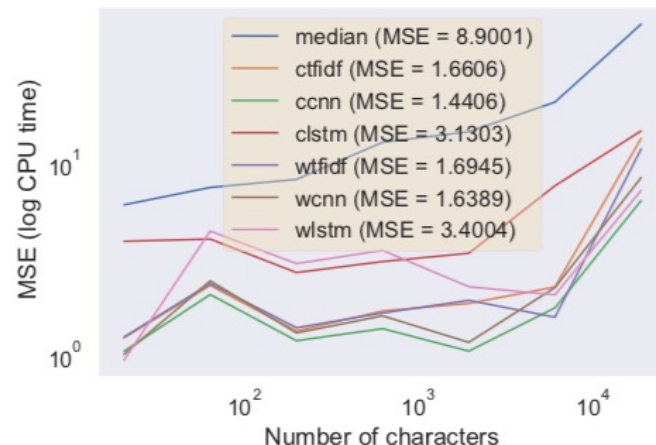
- **Goal:** Predict logarithm of CPU time
- Measure mean squared error (**MSE**)
- Lowest MSE obtained by **character-level models**
 - ccnn is a shallow character-level cnn
- MSE of these models increases as statement complexity (Number of Characters) increases




Results: CPU Time Prediction in Different Settings



(a) Homogeneous Instance



(b) Homogeneous Schema



(c) Heterogeneous Schema

- From left to right, the range of MSE values increases as the problem setting complexity increases
- In each figure, the MSE of models increases as statement complexity increases
- Character-level models obtain lowest MSE and test loss value

Results: Answer Size Prediction

- **Goal:** predict answer size
- Report qerrors in different percentiles of the test data
- **qerror:** shows the factor by which a prediction differs from its true value

Highlights:

- For 50% of queries, it is easy to predict and for top 10% prediction is very difficult
- NN models outperform traditional models which have fixed features
- Character-levels obtain the lowest qerror

Answer size prediction qerror in **SDSS**

Model	50%	75%	80%	85%	90%	95%
median	1	36	50	144	1885	50000
ctfidf	1.13	4.86	10	25	88	727
ccnn	1.36	2.60	3.75	6.79	18	174
clstm	1.07	2.38	3.50	6.79	19	172
wtfidf	1.00	5.37	11.04	31.98	100	879
wcnn	1.33	3.42	5.14	10.93	36	295
wlstm	1.12	2.62	4.27	10.43	30	292



Selected Related Work

- Query Performance Prediction

- [LGMB15] Leis, V., Gubichev, A., Mirchev, A., Boncz, P., Kemper, A., & Neumann, T. (2015). How good are query optimizers, really?. Proceedings of the VLDB Endowment, 9(3), 204-215.
- [LKNC12Li] Jiexing, Arnd Christian König, Vivek Narasayya, and Surajit Chaudhuri. "Robust estimation of resource consumption for sql queries using statistical techniques." Proceedings of the VLDB Endowment 5, no. 11 (2012):
- [BDM19] Bailu Ding, Sudipto Das, Ryan Marcus, Wentao Wu, Surajit Chaudhuri, and Vivek Narasayya. 2019. AI Meets AI: Leveraging Query Executions to Improve Index Recommendations. In Proceedings of the 2019 ACM SIGMOD International Conference on Management of data. SIGMOD'19.

- SQL Query Workloads

- [RTS14] M Jordan Raddick, Ani R Thakar, Alexander S Szalay, and Rafael DC Santos. 2014. Ten Years of SkyServer I: Tracking Web and SQL e- Science Usage. Computing in Science & Engineering 16, 4 (2014), 22–31.
- [JMH16] Shrainik Jain, Dominik Moritz, Daniel Halperin, Bill Howe, and Ed Lazowska. 2016. Sqlshare: Results from a multi-year sql-as-a-service experiment. In Proceedings of the 2016 International Conference on Management of Data. ACM, 281–293.



Contributions: Facilitating SQL Query Composition and Analysis

- Introduce and address 4 problems for predicting **query performance properties - prior to execution**
- Approach is based on using **large-scale real-world query workloads**
- Conduct extensive **workload analysis**
- Adapt data-driven machine learning models
 - **Establish baselines** and assess **feasibility**
- Results show character level models (e.g., ccnn) generalize better under different problem settings

